

MovieLens Recommender System Capstone Project - Report

Sanjana.B

Contents

1 Executive Summary.....	1
2 Method and Analysis.....	2
2.1 Initial data Exploration	2
2.2 Dataset Pre-Processing and Feature Engineering	3
2.3 Rating Distribution	4
2.4 Genre Analysis	14
3 Analysis - Model Building, Evaluation and Insights.....	18
3.1 Naive Baseline Model	18
3.2 Movie-Based Model, a Content-based Approach	18
3.3 Movie + User Model, a User-based approach	18
4 Results.....	19
5 Conclusion.....	20
6 Appendix.....	20
6.1 1a – Project code.	20

1 Executive Summary

The purpose for this project is creating a movie recommendation system using MovieLens dataset.

The version of movielens dataset used for this final assignment contains approximately 10 Millions of movies ratings, divided in 8.5 Millions for training and 1.5 Million for validation or testing. It is a small subset of a much larger dataset with several millions of ratings. Into the training dataset there are approximately

70.000 users and **11.000 different movies** divided in 20 genres such as Action, Adventure, Horror, Drama, Thriller and more.

After a initial data exploration, the recommender systems built on this dataset are evaluated and chosen

based on the RMSE - Root Mean Squared Error that should be at least lower than **0.8649**.

For accomplishing this goal, the **Movie+User Model** is capable to reach a RMSE of **0.863**, that is really good.

2 Method and Analysis

2.1 Initial data Exploration

The 10 Millions dataset is divided into two dataset: edx for training purpose and validation for the validation phase.

The edx dataset contains approximately 9 Millions of rows with 70.000 different users and 11.000 movies with rating score between 0.5 and 5. There is no missing values (0 or NA).

Users	Movies
<int>	<int>
1 69878	10677

The features/variables/columns in both datasets are six:

- **userId** (class integer) Unique user identification number .
- **movieId** (class numeric) Unique movie identification number
- **rating** (class numeric) One user's rating of one movie. Ratings are given out of a 5(1,1.5,2...)
- **timestamp** (class integer) Timestamp for one particular rating by one particular user.
- **title** (class character) Title of each movie with the release year.
- **genres** (class character) list of genres of each movie separated by pipe(|).

First 6 Rows of edx dataset

movieId	title	year	genres
1 31	Dangerous Minds	1995	Drama
2 1029	Dumbo	1941	Animation Children Drama Musical
3 1061	Sleepers	1996	Thriller
4 1129	Escape from New York	1981	Action Adventure Sci-Fi Thriller
5 1172	Cinema Paradiso (Nuovo cinema Paradiso)	1989	Drama
6 1263	Deer Hunter	1978	Drama War

	userId	rating	timestamp
1	1	2.5	1260759144
2	1	3.0	1260759179
3	1	3.0	1260759182
4	1	2.0	1260759185
5	1	4.0	1260759205
6	1	2.0	1260759151

2.2 Dataset Pre-Processing and Feature Engineering

After some initial data exploration, we see that each movie has many different genres. It's necessary to extract and separate them for better consistency. We also notice that the title contains the year of the movie release which can be made into a separate column in the data frame to be used for further analysis. Finally, we can obtain the year and the month for each rating from the timestamp column.

The pre-processing phase is comprised of the following:

1. Format timestamp to a human readable date
2. Extract the month and the year from the date
3. Extract the release year from the title
4. Separate each genre from the list of genres of each movie. It increases the size of both datasets.

After preprocessing the data, edx dataset looks like this:

Processed edx dataset

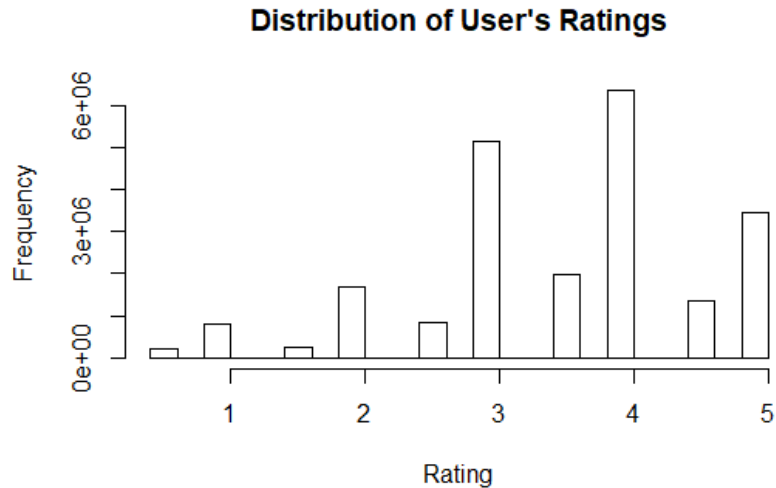
	userId	movieId	rating	title	release	year Rated	month Rated	genre
	<int>	<dbl>	<dbl>	<chr>	<int>	<dbl>	<dbl>	<chr>
1	1	122	5	Boomerang (1992)	1992	96	8	Comedy
2	1	122	5	Boomerang (1992)	1992	96	8	Romance
3	1	185	5	Net, The (1995)	1995	96	8	Action
4	1	185	5	Net, The (1995)	1995	96	8	Crime
5	1	185	5	Net, The (1992)	1995	96	8	Thriller

2.3 Rating frequency Distributions:


Overview of Rating Distribution

According to the given histogram, it can be seen that there are a small number of negative votes. It is noticed that half-Star votes are rarer than “Full-Star” votes.

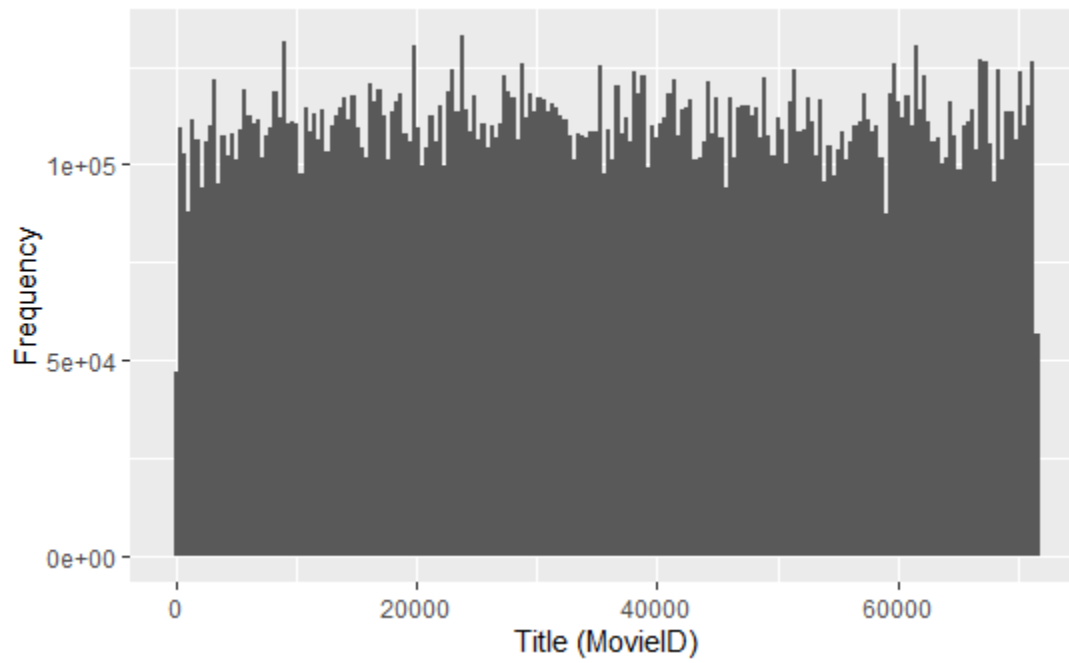
Distribution of User Ratings




Rating frequency Distribution of User Ratings per Title





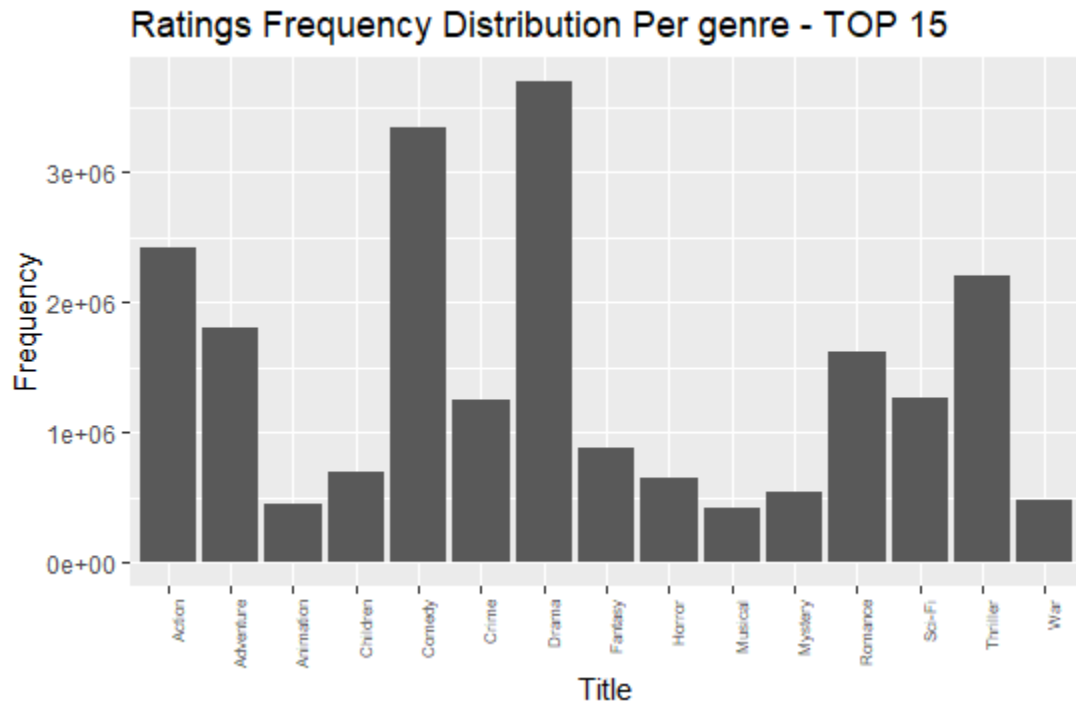
Ratings Frequency Distribution Per Title (MovieID)



Ratings Frequency Distribution







Rating frequency distribution per title(top 10)

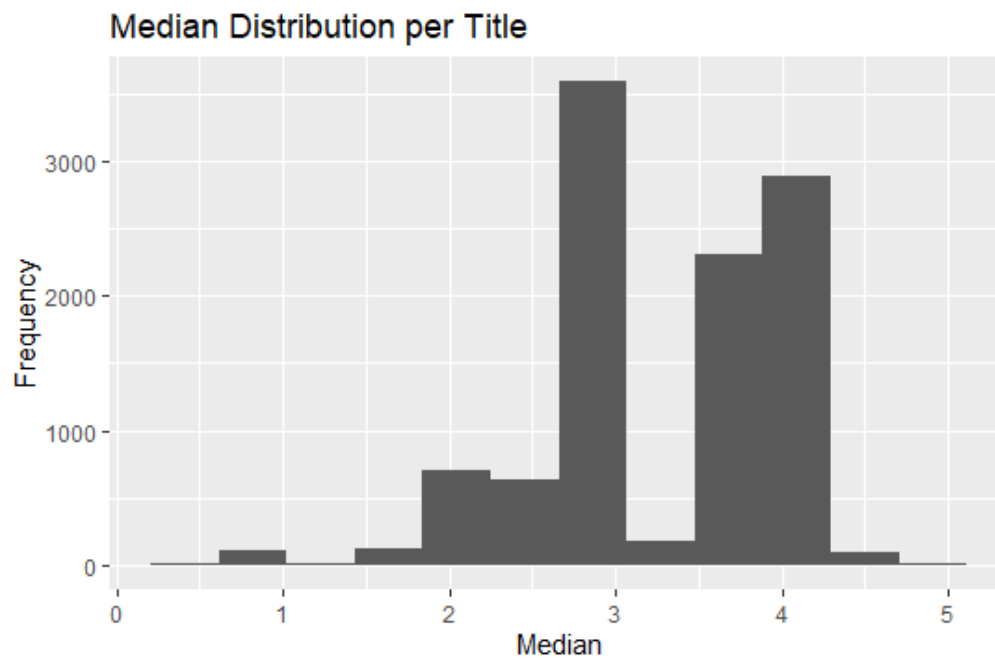
title	count
<chr>	<int>
1 Forrest Gump (1994)	<u>117180</u>
2 Toy Story (1995)	<u>112450</u>
3 Jurassic Park (1993)	<u>111016</u>
4 True Lies (1994)	<u>107790</u>
5 Aladdin (1992)	<u>99880</u>
6 Batman (1989)	<u>91788</u>
7 Pulp Fiction (1994)	<u>89058</u>
8 Lion King, The (1994)	<u>89005</u>
9 Independence Day (a.k.a. ID4) (1996)	<u>88848</u>
10 Silence of the Lambs, The (1991)	<u>85665</u>

Rating frequency distribution per year


release	count
<int>	<int>
1 1995	1968372
2 1994	1635409
3 1996	1475053
4 1999	1094170
5 1997	1074284
6 1998	1025660
7 1993	1025377
8 2000	878117
9 2001	786606
10 2002	671877

Median rating Distributions


Median rating distribution histograms



Median Distribution per Release



Median Distribution per user



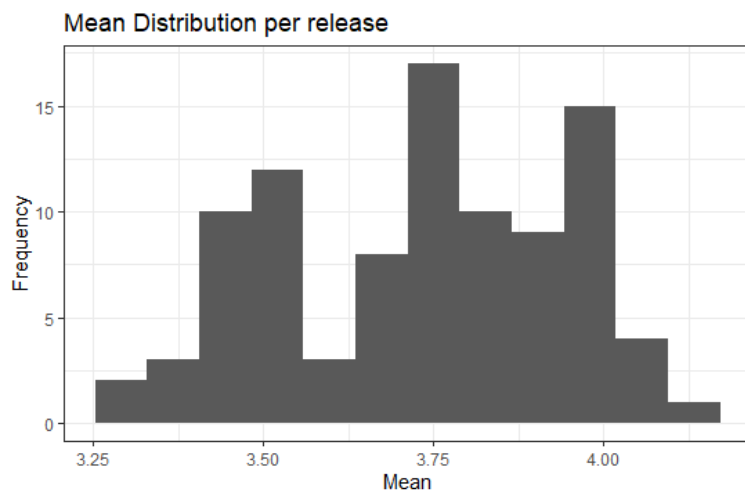
Median rating distribution tables

2.3.1 Median distribution per title(top 10)

title	median
<chr>	<dbl>
1 Blue Light, The (Das Blaue Licht) (1932)	5
2 Class, The (Entre les Murs) (2008)	5
3 Constantine's Sword (2007)	5
4 Fighting Elegy (Kenka erejii) (1966)	5
5 Godfather, The (1972)	5
6 Kids of Survival (1996)	5
7 More (1998)	5
8 Satan's Tango (SÃ;tÃ;ntangÃ³) (1994)	5
9 Shadows of Forgotten Ancestors (1964)	5
10 Shawshank Redemption, The (1994)	5

Median distribution per year(top 10)

release	median
<int>	<dbl>
1 1916	4
2 1918	4
3 1920	4
4 1921	4
5 1922	4
6 1923	4
7 1924	4
8 1925	4
9 1926	4
10 1927	4




Median distribution per user(top 10)

	userId	median
	<int>	<dbl>
1	1	5
2	4	5
3	30	5
4	43	5
5	44	5
6	51	5
7	56	5
8	98	5
9	104	5
10	123	5

Mean Rating distributions

Mean Rating distribution histograms



Mean Rating Distribution Per Release Year- Top 10

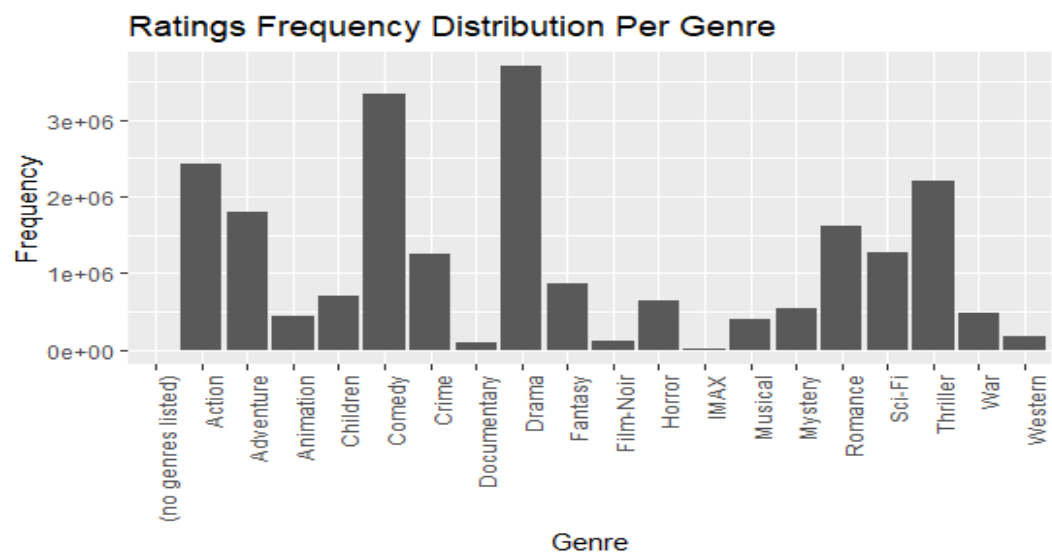
release	mean
<int>	<dbl>
1 1931	4.11
2 1934	4.08
3 1946	4.06
4 1944	4.05
5 1957	4.03
6 1927	4.01
7 1942	4.00
8 1952	4.00
9 1949	4.00
10 1962	3.99

Mean Rating Distribution Per User ID- Top 10


userId	mean
<int>	<dbl>
1 1	5
2 7984	5
3 11884	5
4 13027	5
5 13513	5
6 13524	5
7 15575	5
8 18965	5
9 22045	5
10 26308	5

Genre Analysis

2.4. Rating Distribution per Genre



Rating Distribution per Month



Rating Distribution Tables


Ratings Frequency Distribution Per Genre

genre	count
<chr>	<int>
1 Drama	<u>3693674</u>
2 Comedy	<u>3343505</u>
3 Action	<u>2418271</u>
4 Thriller	<u>2197031</u>
5 Adventure	<u>1802802</u>
6 Romance	<u>1616899</u>
7 Sci-Fi	<u>1267340</u>
8 Crime	<u>1254235</u>
9 Fantasy	<u>874268</u>
10 Children	<u>696914</u>
11 Horror	<u>652597</u>
12 Mystery	<u>537130</u>
13 War	<u>482621</u>
14 Animation	<u>440852</u>
15 Musical	<u>408685</u>
16 Western	<u>178863</u>
17 Film-Noir	<u>112168</u>
18 Documentary	<u>87957</u>
19 IMAX	<u>7739</u>
20 (no genres listed)	5



Ratings Frequency Distribution Per Month

	month Rated	count
	<dbl>	<int>
1	11	2394359
2	12	2203980
3	10	2035272
4	7	1972738
5	1	1843616
6	6	1837634
7	3	1827162
8	8	1773138
9	5	1638183
10	4	1633203
11	2	1515358
12	9	1398913

2.4.2 Mean Distribution per Genre



2.4.3 Median Distribution per Genre



Median distribution per genre

genre	median
<chr>	<dbl>
1 Animation	4
2 Crime	4
3 Documentary	4
4 Drama	4
5 Film-Noir	4
6 IMAX	4
7 Musical	4
8 Mystery	4
9 Romance	4
10 War	4

Mean distribution per genre

genre	mean
<chr>	<dbl>
1 Film-Noir	4.01
2 Documentary	3.78
3 War	3.78
4 IMAX	3.76
5 Mystery	3.68
6 Drama	3.67
7 Crime	3.67
8 Animation	3.60
9 Musical	3.56
10 Western	3.56

3 Analysis - Model Building and Evaluation

3.1 Naive Baseline Model

The simplest model possible, is a Naive Model that predicts the mean for all cases. In this case, the mean is **3.512465**, approximately **3.5**.

3.1.1 Naive Mean-Baseline Model

The formula used is:

$$Y_{u,i} = \hat{\mu} + u_{u,i}$$

With $\hat{\mu}$ is the mean and $\epsilon_{i,u}$ is the independent errors sampled from the same distribution centered at 0.

$\hat{\mu}$ is 3.527021 and the rmse on the validation set is **1.052698**, approximately **1.05**. It is larger than the target RMSE (below 0.87) and that indicates poor performance for the model.

3.2 Movie-Based Model, a Content-based Approach

The first Non-Naive Model takes into account the type of movie. In this case the movies that are rated higher or lower respect to each other.

The formula used is:

$$Y_{u,i} = \hat{\mu} + b_i + \epsilon_{u,i}$$

With $\hat{\mu}$ is the mean and $\epsilon_{i,u}$ is the independent errors sampled from the same distribution centered at 0. The b_i is a measure for the popularity of movie i , i.e. the bias of movie i .

The RMSE on the validation dataset is **0.9417822**. It is better than the Naive Mean-Baseline Model but it is also much higher than the target RMSE (below 0.87) and that indicates poor performance for the model.

3.3 Movie + User Model, a User-based approach

The second Non-Naive Model considers that each user has different preference of movies and rate differently according to their perspectives.

The formula used is:

$$Y_{u,i} = \hat{\mu} + b_i + b_u + \epsilon_{u,i}$$

With $\hat{\mu}$ is the mean and $\epsilon_{i,u}$ is the independent errors sampled from the same distribution centered at 0. The b_i is a measure for the popularity of movie i , i.e. the bias of movie i . The b_u is a measure for the mildness of user u , i.e. the bias of user u .

The RMSE on the validation dataset is 0.8639665 which is very good. The Movie+User Based Model

has obtained the required performance but by applying regularization, the model can be improved by some amount.

4 Results

This is the summary results for all the model built, trained on edx dataset and validated on the validation dataset.

[rmse_results](#)

model	RMSE
1 Naive Mean-Baseline Model	1.0526979
2 Movie-Based Model	0.9417822
3 Movie+User Based Model	0.8639665

5 Conclusion

After training different models, it's very clear that movieId and userId are the main contributors for prediction. Without regularization, the model has achieved the desired performance, but by applying regularization and adding the genre predictor, performance can be improved and RMSE can be reduced.

6 Appendix

```
##Installing required packages
```

```
install.packages("forcats")
```

```
install.packages("kableExtra")
```

```
#Loading the libraries for the project
```

```
library(ggplot2)
```

```
library(kableExtra)
```

```
library(stringr)
```

```
library(tidyr)
```

```
library(tibble)
```

```
library(tidyverse)
```

```
library(dslabs)
```

```
library(dbplyr)
```

```
library(caret)
```

```
library(broom)
```

```
library(naivebayes)
```

```
library(pdftools)
```

```
library(rvest)
```

```
library(timeDate)
```

```
library(readr)
```

```
library(purrr)
```

```
library(lubridate)
```

```
library(labeling)
```

```
library(dplyr)
```

```
library(e1071)
```

```
library(data.table)
```

```
# Create edx set, validation set, and final file
```

```
dl <- tempfile()
```

```
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)
```

```

ratings <- fread(text = gsub("::", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),
  col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\::", 3)
colnames(movies) <- c("movieId", "title", "genres")
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],
  title = as.character(title),
  genres = as.character(genres))

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind="Rounding")
# if using R 3.5 or earlier, use `set.seed(1)` instead
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

# Exploratory Data Analysis

## Initial data Exploration

#The 10 Millions dataset is divided into two dataset: "edx" for training purpose and "validation"
for the validation phase.

***edx dataset**
#Defining the root mean square error function

```

```

RMSE <- function(true_ratings = NULL, predicted_ratings = NULL) {
  +   sqrt(mean((true_ratings - predicted_ratings)^2))}

#Viewing first 6 rows of edx and validation dataframes
head(edx)
head(validation)

#Finding the number of unique users and movies in the edx dataframe
edx %>% summarize(Users = n_distinct(userId), Movies = n_distinct(movieId))

#*****Dataset Pre-Processing and Feature Engineering*****

# Convert timestamp to a human readable date

#Adding a date column to both frames with time zone as Greenwich Mean Time
validation$date <- as.POSIXct(validation$timestamp, origin='1970-01-01',tz="GMT")
edx$date <- as.POSIXct(edx$timestamp,origin='1970-01-01',tz="GMT")

#Viewing first 6 rows of edx and validation dataframes
head(validation)
head(edx)

#Adding month and year columns to both dataframes
validation$month_rated <- format(validation$date,"%m")
validation$year_rated <- format(validation$date,"%y")
edx$year_rated <- format(edx$date,"%y")
edx$month_rated <- format(edx$date,"%m")

#validation <- validation %>% mutate(title = str_squish(title)) %>% extract(title,c("titleTemp",
"release"),regex = "^(.*) \\([([0-9 \\-]*)\\)$",remove = F) %>% mutate(release =
if_else(str_length(release) > 4,as.integer(str_split(release, "-",simplify =
T)[1]),as.integer(release)))

#Adding release date column to validation
validation <- validation %>%mutate(title = str_trim(title)) %>%extract(title, c("titleTemp",
"release"),regex = "^(.*) \\([([0-9 \\-]*)\\)$",remove = F) %>% mutate(release =
if_else(str_length(release) > 4,as.numeric(str_split(release, "-",simplify =
T)[1]),as.numeric(release))) %>% mutate(title = if_else(is.na(titleTemp),title,titleTemp))

```

```
validation<-validation %>% select(-titleTemp)
```

```
#Viewing first 6 rows of edx and validation data frames
```

```
head(validation)
```

```
head(edx)
```

```
#Adding release date column to edx
```

```
edx <- edx %>% mutate(title = str_squish(title)) %>% extract(title,c("titleTemp",  
"release"),regex = "^(.*) \\([([0-9 \\-]*)\\)$",remove = F) %>% mutate(release =  
if_else(str_length(release) > 4,as.integer(str_split(release, "-",simplify =  
T)[1]),as.integer(release)))
```

```
edx<-edx %>% select(-titleTemp)
```

```
head(edx)
```

```
# Extract the genre in validation datasets
```

```
validation <- validation %>%mutate(genre = fct_explicit_na(genres,na_level = "(no genres  
listed)")) %>%separate_rows(genre, sep = "\\|")
```

```
# Extract the genre in edx datasets
```

```
edx <- edx %>% mutate(genre = fct_explicit_na(genres,na_level = "(no genres listed)")) %>%  
separate_rows(genre,sep = "\\|")
```

```
#Viewing first 6 rows of edx and validation dataframes
```

```
head(edx)
```

```
head(validation)
```

```
#Removing the unnecessary columns in both data frames
```

```
validation <- validation %>% select(-genres,-timestamp,-date)
```

```
edx <- edx %>% select(-genres,-timestamp,-date)
```

```
# Convert the columns into the desired data type
```

```
validation$month Rated <- as.numeric(validation$month Rated)
```

```
validation$year Rated <- as.numeric(validation$year Rated)
```

```
edx$year Rated <- as.numeric(edx$year Rated)
```

```
edx$month Rated <- as.numeric(edx$month Rated)
```

```

#####Processed edx dataset#####

head(edx)

#####Processed validation dataset#####

head(validation)

#####END OF DATA PROCESSING#####

#####MOVIE AND RATING HISTOGRAMS#####

#movies released per year
hist(edx$release)

#rating distribution
hist(edx$rating, main="Distribution of User's Ratings", xlab="Rating")

#####RATING FREQUENCY DISTRIBUTION HISTOGRAMS#####

### Numbers of Ratings per Movie

ggplot(edx, aes(movieId)) + theme_grey() + geom_histogram(bins=700) + labs(title = "Ratings
Frequency Distribution Per Title (MovieID)",x = "Title (MovieID)",y = "Frequency")

#Number of ratings per user id

ggplot(edx, aes(userId)) +theme_grey() +geom_histogram(bins=200) +labs(title = "Ratings
Frequency Distribution Per Title (MovieID)",x = "Title (MovieID)",y = "Frequency")

#####RATING FREQUENCY DISTRIBUTION PLOTS#####

#Rating frequency distribution

ggplot(edx, aes(rating)) + theme_grey() + geom_histogram() +labs(title = "Ratings Frequency
Distribution ", x = "Rating", y = "Frequency")

```


#Ratings Frequency Distribution Per Title - TOP 20 Movies

```
edx %>% group_by(title) %>% summarise(count = n()) %>% arrange(desc(count))
%>% head(n=20) %>% ggplot(aes(title, count)) +theme_gray() +geom_col() +theme(axis.text.x
= element_text(angle = 90, hjust = 1, size = 7)) +labs(title = "Ratings Frequency Distribution Per
Title - TOP 20 Movies",x = "Title",y = "Frequency")
```

#Ratings Frequency Distribution Per Year - TOP 10

```
edx %>% group_by(release) %>% summarise(count = n()) %>% arrange(desc(count))
%>% head(n=10) %>% ggplot(aes(release, count)) +theme_gray() +geom_col()
+theme(axis.text.x = element_text(angle = 90, hjust = 1, size = 6)) +labs(title = "Ratings
Frequency Distribution Per Year - TOP 10",x = "Title", y = "Frequency")
```

#Ratings Frequency Distribution Per genre - TOP 15

```
edx %>% group_by(genre) %>% summarise(count = n()) %>% arrange(desc(count)) %>%
head(n=15) %>% ggplot(aes(genre, count)) + theme_gray() + geom_col() +
theme(axis.text.x = element_text(angle = 90, hjust = 1, size = 6)) +
labs(title = "Ratings Frequency Distribution Per genre - TOP 15", x = "Title", y = "Frequency")
head(edx)
```

*****RATING FREQUENCY DISTRIBUTION TABLES*****

#Rating frequency distribution per title

```
edx %>% group_by(title) %>% summarise(count = n()) %>% arrange(desc(count)) %>%
head(n=25)
```

#Rating frequency distribution per release year

```
edx %>% group_by(release) %>% summarise(count = n()) %>% arrange(desc(count)) %>%
head(n=25)
```

*****MEDIAN DISTRIBUTION HISTOGRAMS*****

Median Distribution per Title (Movie) histogram

```
edx %>%
  group_by(title) %>%
  summarise(median = median(rating)) %>%
  ggplot(aes(median)) +
  theme_gray() +
  geom_histogram(bins=12) +
  labs(title = "Median Distribution per Title",x = "Median",y = "Frequency")
```

Median distribution per release (year) histogram

```
edx %>%
  group_by(release) %>%
  summarise(median = median(rating)) %>%
  ggplot(aes(median)) +
  theme_gray() +
  geom_histogram(bins=12) +
  labs(title = "Median Distribution per Release",x = "Median",y = "Frequency")
```

Median distribution per user histogram

```
edx %>%
  group_by(userId) %>%
  summarise(median = median(rating)) %>%
  ggplot(aes(median)) +
  theme_gray() +
  geom_histogram(bins=12) +
  labs(title = "Median Distribution per user",x = "Median",y = "Frequency")
```

```
#####MEDIAN DISTRIBUTION TABLES#####
```

```
### Median distribution per title(movie) table
```

```
edx %>%  
  group_by(title) %>%  
  summarise(median = median(rating)) %>%  
  arrange(desc(median)) %>%  
  head(n=25)
```

```
### Median distribution per release year table
```

```
edx %>%  
  group_by(release) %>%  
  summarise(median = median(rating)) %>%  
  arrange(desc(median)) %>%  
  head(n=25)
```

```
### Median distribution per user table
```

```
edx %>%  
  group_by(userId) %>%  
  summarise(median = median(rating)) %>%  
  arrange(desc(median)) %>%  
  head(n=25)
```

```
#####MEAN DISTRIBUTION HISTOGRAMS#####
```

```
###Mean distribution per title histogram
```

```
edx %>%  
  group_by(title) %>%  
  summarise(mean = mean(rating)) %>%  
  ggplot(aes(mean)) +  
  theme_light() +  
  geom_histogram(bins=12) +  
  labs(title = "Mean Distribution per Title",x = "Mean",y = "Frequency")
```

```
###Mean distribution per release histogram
```

```
edx %>%  
  group_by(release) %>%  
  summarise(mean = mean(rating)) %>%  
  ggplot(aes(mean)) +  
  theme_bw() +  
  geom_histogram(bins=12) +  
  labs(title = "Mean Distribution per release",x = "Mean", y = "Frequency")
```

```
###Mean distribution per user histogram
```

```
edx %>%  
  group_by(userId) %>%  
  summarise(mean = mean(rating)) %>%  
  ggplot(aes(mean)) +  
  theme_bw() +  
  geom_histogram(bins=12) +  
  labs(title = "Mean Distribution per user",x = "Mean",y = "Frequency")
```

```
#####MEAN DISTRIBUTION TABLES#####
```

```
##MEAN DISTRIBUTION PER TITLE TABLE
```

```
edx %>%  
  group_by(title) %>%  
  summarise(mean = mean(rating)) %>%  
  arrange(desc(mean)) %>%  
  head(n=10)
```

```
##MEAN DISTRIBUTION PER RELEASE YEAR TABLE
```

```
edx %>%  
  group_by(release) %>%  
  summarise(mean = mean(rating)) %>%  
  arrange(desc(mean)) %>%  
  head(n=10)
```

```
##MEAN DISTRIBUTION PER USER TABLE
```

```
edx %>%  
  group_by(userId) %>%  
  summarise(mean = mean(rating)) %>%  
  arrange(desc(mean)) %>%  
  head(n=10)
```

```
##*****Rating distributions*****
```

```
### Rating Distribution per Genre
```

```
##*Overview of Rating distribution over Genre**
```

```
edx %>%  
  group_by(genre) %>%  
  summarise(count = n()) %>%  
  ggplot(aes(genre, count)) +  
  theme_gray() +  
  geom_col() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Ratings Frequency Distribution Per Genre", x = "Genre", y = "Frequency")
```

```
##*Overview of Rating distribution over months**
```

```
edx %>%  
  group_by(month_rated) %>%  
  summarise(count = n()) %>%  
  ggplot(aes(month_rated, count)) +  
  theme_gray() +  
  geom_col() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Ratings Frequency Distribution Per month", x = "month", y = "Frequency")
```

```
#####RATING FREQUENCY TABLES#####
```

```
#rating per genre
```

```
edx %>%  
  group_by(genre) %>%  
  summarise(count = n()) %>%  
  arrange(desc(count))
```

```
#rating per month
```

```
edx %>%  
  group_by(month_rated) %>%  
  summarise(count = n()) %>%  
  arrange(desc(count))
```

```
#####Rating distribution plots#####
```

```
### Mean Distribution per Genre
```

```
edx %>%  
  group_by(genre) %>%  
  summarise(mean = mean(rating)) %>%  
  ggplot(aes(genre, mean)) +  
  theme_bw() +  
  geom_col() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Mean Distribution per Genre", x = "Genre", y = "Mean")
```

```
### Median Distribution per Genre
```

```
edx %>%  
  group_by(genre) %>%  
  summarise(median = median(rating)) %>%  
  ggplot(aes(genre, median)) +  
  theme_grey() +  
  geom_col() +  
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +  
  labs(title = "Median Distribution per Genre", x = "Genre", y = "Median")
```

```

### Mode Distribution per Genre
edx %>%
  group_by(genre) %>%
  summarise(mode = mode(rating)) %>%
  ggplot(aes(genre, mode)) +
  theme_gray() +
  geom_col() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  labs(title = "Mode Distribution per Genre", x = "Genre", y = "Mode")

#####Distribution tables#####

####median distribution per genre

edx %>%

  group_by(genre) %>%
  summarise(median = median(rating)) %>%
  arrange(desc(median)) %>%
  head(n=10)

####mean distribution per genre

edx %>%
  group_by(genre) %>%
  summarise(mean = mean(rating)) %>%
  arrange(desc(mean)) %>%
  head(n=10)

##### Analysis - Model Building and Evaluation#####

## Naive Baseline Model
mean(edx$rating)

### Naive Mean-Baseline Model

# Calculate the average of all movies
mu_hat <- mean(edx$rating)
mu_hat

```

```

# Predict the RMSE on the validation set
rmse_mean_result <- RMSE(validation$rating, mu_hat)
rmse_mean_result

# Creating a results dataframe that contains all RMSE results
rmse_results <- data.frame(model="Naive Mean-Baseline Model", RMSE=rmse_mean_result)
rmse_results

## Movie-Based Model, a Content-based Approach
# Calculate the average per each movie

movie_rmse <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = mean(rating - mu_hat))
movie_rmse

# Predict the ratings for validation dataset

rmse_validation <- validation %>%
  left_join(movie_rmse, by='movieId') %>%
  mutate(pred = mu_hat + b_i) %>%
  pull(pred)
rmse_validation

rmse_validation_result <- RMSE(validation$rating, rmse_validation)
rmse_validation_result

# Adding the results to the results dataset
rmse_results <- rmse_results %>% add_row(model="Movie-Based Model",
RMSE=rmse_validation_result)
rmse_results

#The RMSE on the ``validation`` dataset is 0.9417822. It is slightly better than the Naive
Mean-Baseline Model, but it is also far from the required RMSE (below 0.87) leading to poor
performance for the model.

```



```
## Movie + User Model, a User-based approach
```

```
# Calculate the average by movie
```

```
movie_avgs <- edx %>%  
  group_by(movieId) %>%  
  summarize(b_i = mean(rating - mu_hat))
```

```
# Calculate the average for every user
```

```
user_avgs <- edx %>%  
  left_join(movie_avgs, by='movieId') %>%  
  group_by(userId) %>%  
  summarize(b_u = mean(rating - mu_hat - b_i))
```

```
# Compute the predicted ratings on validation dataset
```

```
rmse_movie_plus_user_model <- validation %>%  
  left_join(movie_avgs, by='movieId') %>%  
  left_join(user_avgs, by='userId') %>%  
  mutate(pred = mu_hat + b_i + b_u) %>%  
  pull(pred)
```

```
rmse_movie_plus_user_model_result<-RMSE(validation$rating, rmse_movie_plus_user_model)
```

```
# Adding the results to the results dataset
```

```
rmse_results <- rmse_results %>% add_row(model="Movie+User Based Model",  
  RMSE=rmse_movie_plus_user_model_result)
```

```
rmse_results
```

```
#The movie plus user based model has achieved the
```

```
#required resultant RMSE of 0.863 which is less than 0.8649
```

```
#The model can be improved by using regularisation technique.
```